

Lab session 3

Machine Learning for Behavioral Data (CS-421)

March 10, 2021


Today

- **08:15 08:30 Part I Introduction**
 - Feedback from tutorial 2
 - Introduction to Scikit-learn & SpeakUp
 - **08:30 09:10 Part II Tutorial on Regression**
 - **09:10 09:15 BREAK**
 - **09:15 09:50 Part III Tutorial on Classification**
 - **09:50 10:00 Part IV Homeworks and Next Steps**
 - Debrief on homework 1
 - Introduction to homework 2
 - Closing and feedback
-

Feedback last tutorial (n=2)


1

The volume and complexity of the material were appropriate, from 1 (totally disagree) to 5 (totally agree)



2

The lab session was well-paced and clear, from 1 (totally disagree) to 5 (totally agree)



3

Theory and hands-on activities were well balanced, from 1 (totally disagree) to 5 (totally agree)



Feedback last tutorial (n=1)

4

What do you like about this lab? What should be kept?

Respondent

Response

I really like the fact that you take the time to walkthrough things and ask about our background. Also it's really cool that you provide us with useful and clear ressources like the cheatsheets.

Total responses to
question

1/2

Other points from previous feedback forms:

- Keep SpeakUp and little quizzes.
- Detailed explanations and walkthrough guidance.
- Openness and time to answer questions.
- Well-documented notebooks.

Feedback last tutorial (n=1)

5

What do you dislike about this lab? What should be changed (and how)?

Respondent

Response

Maybe sometimes you can go a little faster for certain things but it's ok

Other points from previous feedback forms:

- Hard to know what library or function to use and when. Provide supporting information.
- Provide a bit more time to complete exercise cells with nice results.
- Control the pace of the tutorial and provide more non-copy-and-pasting exercises.

Feedback last tutorial (n=2)

8

Any additional comment?

Respondent

Response

I think maybe there should be more time for answering questions on the lab or homework. But I do like the tutorials and speakup so it's not easy...maybe balance more and have less homeworks/labs

Would it be possible to record the lab sessions ? Also lecture 2 is still not uploaded and I couldn't attend it on Monday.

Total responses to
question

2/2

Other points from previous feedback forms:

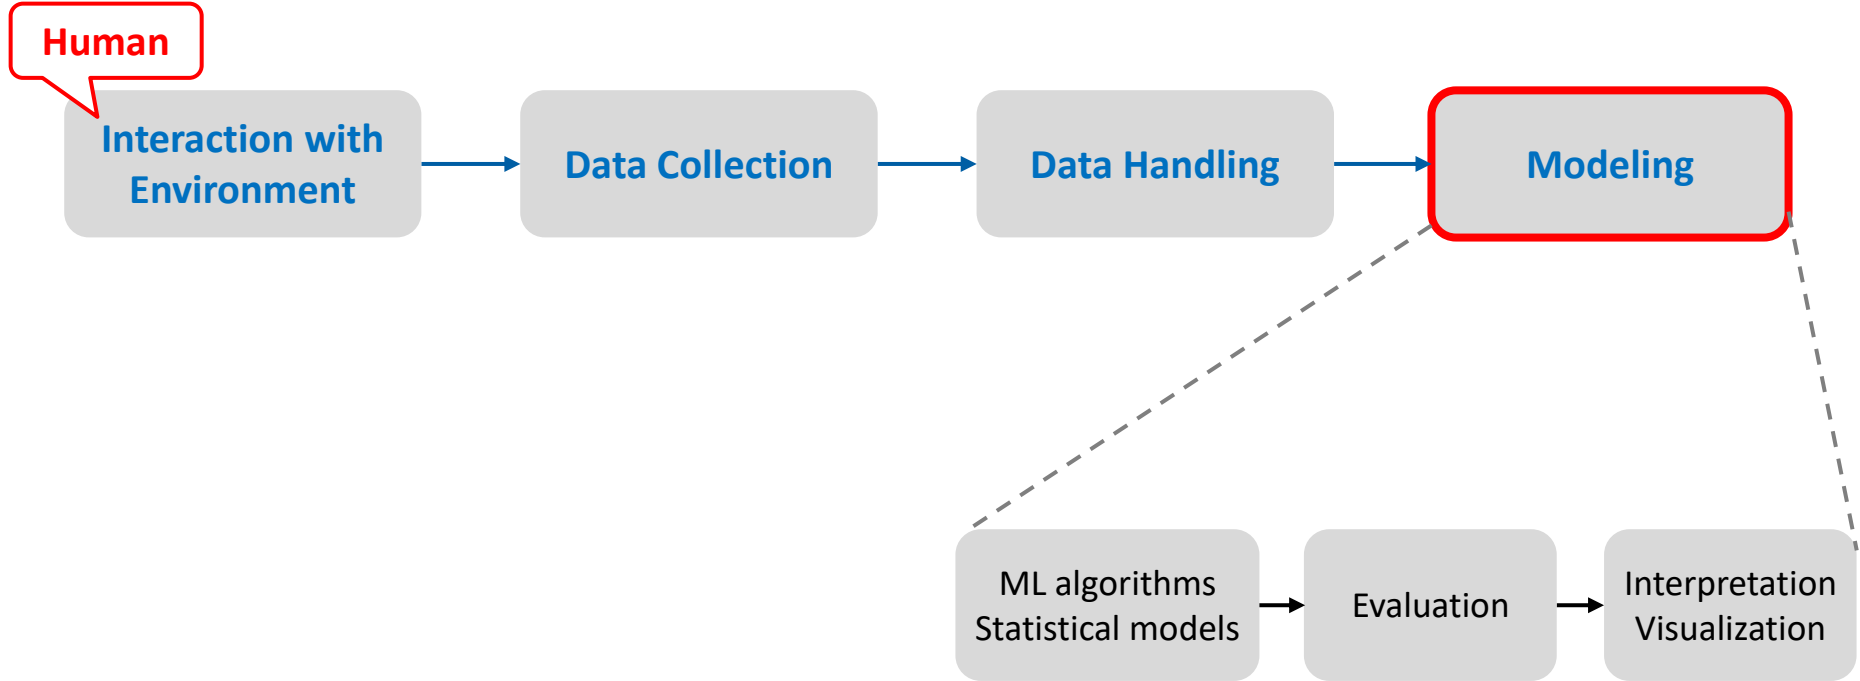
- Time of uploading for the material.
- Investigating the possibility of recording lab sessions.

Where we are

Week	Lecture	Lab Sessions	Project
1	Introduction	Tutorial	
2	Data Handling	Tutorial + Homework	
3	Regression & Classification	Tutorial + Homework	
4	Model Selection & Evaluation	Tutorial + Homework	Presentation of data sets and research questions
5	Latent Variable Models	Tutorial + Homework	M1: Preferences on team members and data sets
6	Unsupervised Learning	Tutorial + Homework + PO	
7	Spring Break	Spring Break	Spring Break

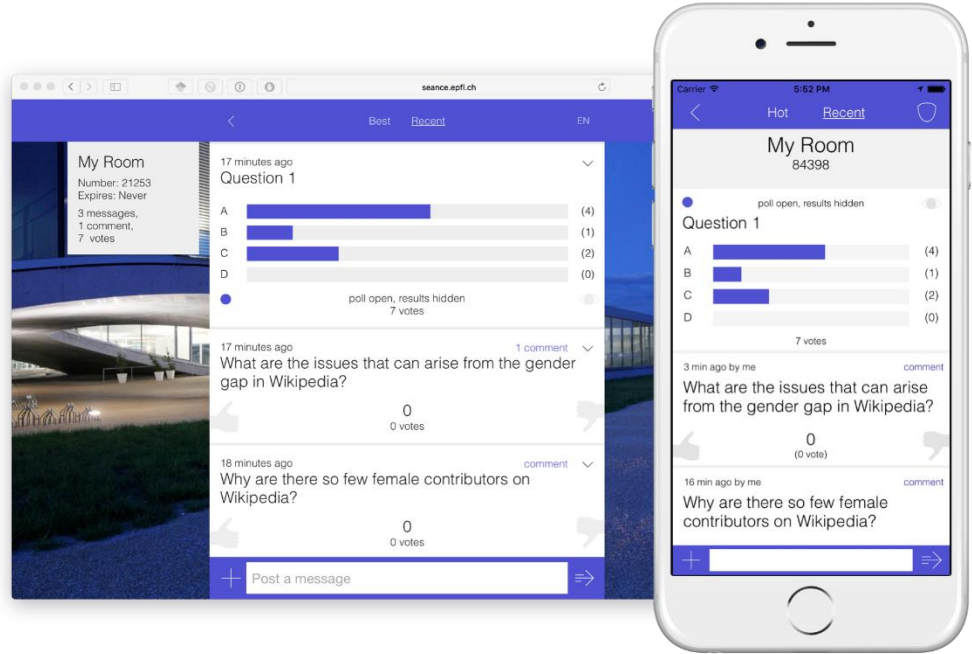
PO = project office hours

ML for Behavioral Data: Modeling



SpeakUp

- **Android / iOS:**
<http://speakup.info/>
- **Web App:**
<https://web.speakup.info/>
- **Room number: 95305**



Scikit-learn

SpeakUp #1: How do you feel about scikit-learn?

A: I've never heard of scikit-learn.

B: I'm **not confident at all** about using scikit-learn.

C: I'm **slightly confident** about using scikit-learn.

D: I'm **fairly confident** about using scikit-learn.

E: I'm **very confident** about using scikit-learn.



Quiz

SpeakUp #2:

What is the default score in scikit-learn, when using a classifier?

A: Recall

B: ROC-AUC

✓ C: Accuracy

D: Precision

Quiz

Speakup #3:

Which classifier's method should you call to get the class predictions?

A: `predict(X, y)`

✓ B: `predict(X)`

C: `predict_proba(X, y)`

D: `predict_proba(X)`

Quiz

Speakup #4:

Which command is used to train a model in scikit-learn?

A: train

B: fit_transform

✓ C: fit

D: transform

Quiz

Speakup #5:

Which line of code returns X with values scaled between 0 and 1?

A: `MinMaxScaler().fit(X)`

B: `MinMaxScaler().fit(X, y)`

C: `StandardScaler().fit_transform(X)`

✓ D: `MinMaxScaler().fit_transform(X)`

Quiz

Speakup #6: Given an array X with 10,000 observations and 50 features and an array y with the label for each observation, supposed that we run the command:

$X_{train}, X_{test}, y_{train}, y_{test} = \text{train_test_split}(X, y, \text{test_size}=.25)$

What would be the dimensions of X_{train} , y_{train} , X_{test} , and y_{test} ?

- ✓ A: $X_{train}:(7500,50)$ $y_{train}:(7500,)$ $X_{test}:(2500,50)$ $y_{test}(2500,)$
 - B: $X_{train}:(7500,50)$ $y_{train}:(2500,2)$ $X_{test}:(7500,50)$ $y_{test}(2500,2)$
 - C: $X_{train}:(10000,50)$ $y_{train}:(10000,)$ $X_{test}:(10000,50)$ $y_{test}(10000,)$
 - D: $X_{train}:(2500,)$ $y_{train}:(2500,50)$ $X_{test}:(7500,)$ $y_{test}(7500,50)$
-

Quiz

Speakup #7: Suppose that we have run the line of code: *data = load_iris()*. Identify the line of code below where an error will occur.

A: `X_train, X_test, y_train, y_test = train_test_split(data['data'], data['target'])`

B: `knn = KNeighborsClassifier(n_neighbors=1)`

C: `knn.fit(X_train, y_train)`

✓ D: `y_pred = knn.predict(X_test, y_test)`

scikit-learn
algorithm cheat-sheet

START

get more data

>50 samples

classification

kernel approximation

SGD Classifier

NOT WORKING

KNeighbors Classifier

NOT WORKING

SVC

Ensemble Classifiers

NOT WORKING

Text Data

Naive Bayes

YES

<100K samples

Linear SVC

NOT WORKING

regression

SGD Regressor

Lasso

ElasticNet

YES

few features should be important

<100K samples

NOT WORKING

SVR(kernel='rbf')

EnsembleRegressors

NOT WORKING

RidgeRegression

SVR(kernel='linear')

clustering

Spectral Clustering

GMM

NOT WORKING

KMeans

<10K samples

number of categories known

YES

YES

MiniBatch KMeans

NO

<10K samples

MeanShift

VBGM

NO

tough luck

dimensionality reduction

Randomized PCA

NOT WORKING

<10K samples

YES

Isomap

Spectral Embedding

NOT WORKING

LLE

kernel approximation

predicting a category

YES

do you have labeled data

YES

predicting a quantity

YES

just looking

YES

predicting structure

NO

Back

scikit-learn


learn.org/stable/tutorial/machine_learning_map/index.html

Scikit-learn resources

Choose from below according to your **level**:


- https://www.tutorialspoint.com/scikit_learn/index.htm
 - Walkthrough explanation (text + code)
 - From modelling to regressions, from classification to clustering
 - <https://scikit-learn.org/stable/tutorial/index.html>
 - Extended collection of official tutorials
 - More examples with several datasets.
 - https://s3.amazonaws.com/assets.datacamp.com/blog_assets/Scikit_Learn_Cheat_Sheet_Python.pdf
 - Scikit-learn cheat sheet on one A4 page
-

Further readings



Géron, A. (2019). **Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow: Concepts, tools, and techniques to build intelligent systems.** O'Reilly Media.

Hackeling, G. (2017). **Mastering Machine Learning with scikit-learn.** Packt Publishing.



Questions?



Tutorial 3.1 Agenda

- **10 mins** Fundamentals on Regression
 - Data generation, data split, data scaling, model development, model predictions, model effectiveness
- **20 mins** Hands-on Use Case on Regression
- **10 mins** Final Discussion

Tutorial 3.1 Hands on

We will use Noto, but feel free to use your own environment:

- Go to <https://noto.epfl.ch/>
- Login with your GASPAR
- If you have **NOT** already cloned the repository:
 - Go to Git → Clone → <https://github.com/d-vet-ml4ed/mlbd>
- Go to Git → Pull
- Go through [Tutorials/Tutorial03/Notebooks/Regression.ipynb](#)

Discussion and questions



Tutorial 3.2 Agenda

- **5 mins** Fundamentals on Classification
 - Data generation, data split, data scaling, model development, model predictions, model effectiveness
- **20 mins** Hands-on Use Case on Classification
- **10 mins** Final Discussion

Tutorial 3.2 Hands on

We will use Noto, but feel free to use your own environment:

- Go to <https://noto.epfl.ch/>
- Login with your GASPAR
- If you have **NOT** already cloned the repository:
 - Go to Git → Clone → <https://github.com/d-vet-ml4ed/mlbd>
- Go to Git → Pull
- Go through [Tutorials/Tutorial03/Notebooks/Classification.ipynb](#)

Discussion and questions



Debrief last homework

Homework 1: Data Handling

Introduction

In this homework, you will apply different data exploration, cleaning, and visualization techniques. It is very important to take some time to understand the data.

The homework is due **Mar 09, 2021 23:59 CET**. The fully-run notebook must be uploaded to your private GitHub private repository. We will only grade the cells with the heading

GRADED CELL

If you have any questions, feel free to use the Q&A forum in Moodle.

7 required exercises will be graded:

Section	Part	Required Function	Points
1	Data Exploration	get_feature_stats	20
2	Data Cleaning	handle_missing_values	20
2	Data Cleaning	handle_inconsistent_data	20
2	Data Cleaning	handle_skewness	20
3	Visualization	plot_correlation	10
3	Visualization	plot_grades	5
3	Visualization	Written interpretation	5
Total Points			100

The solution of this homework will be uploaded on the public GitHub repository at:

https://github.com/d-vet-ml4ed/mlbd/tree/main/Homeworks/Homework01/Homework01-DataHandling_sol.ipynb

Questions?



Introduction to homework 2

The link to the homework will be available slightly after this lab session.

We will make an announcement in the forum, when the homework will be published.

Complete the exercises in the .ipynb notebook, commit your work, and **push the changes to GitHub**. TAs will pull the completed assignment after the deadline **Mar 16, 2021 23:59 CET** (any changes after the deadline will not be considered)

If you have any questions, feel free to use the Q&A forum in Moodle!

GitHub Classroom

GitHub Education



You're ready to go!

Important upcoming dates

- Mar 15, 2021 15:15 17:00
 - Lecture #4
- By Mar 16, 2021 23:59 CET
 - Submission Deadline for Homework #2
- Mar 17, 2021 08:15 10:00
 - Lab Session #4

Feedback for this lab session

<https://moodle.epfl.ch/mod/questionnaire/view.php?id=1133047>



Quick Anonymous Feedback on Lecture 3



Quick Anonymous Feedback on Lab Session 3



Quick Anonymous Feedback on Homework 2



Questions?

